

Key Concepts:

- Insert and Delete from B+ tree (4 marks)
 - Insert with linear hashing (2 marks)
 - Perform RAID and index calculations (6 marks)
 - Code an iterator in Java (lab 4) (5 marks)
 - Create relational query plans in Java (5 marks)

1GB = 10^9 bytes

Topics:

I) Storage

a. Memory performance calculations

- i. Transfer time = $\text{memoryTransferSize} / \text{bandwidth}$
- ii. Read time = $\text{readDataSize} / \text{readBandwidth}$
- iii. Write bandwidth = $\text{IOPS} * \text{writeRequests}$

b. Storing records in memory

- i. A record consists of one or more fields grouped together
 1. Each Tuple of a relation is a record
- ii. Two main types of records:
 1. Variable length: size of the record varies
 2. Fixed length: all records are the same size

c. Variable formats

- i. Useful cases:
 1. the data does not have a regular structure in most cases
 2. the data values are sparse in the records
 3. there are repeating fields in the records
 4. the data evolves quickly so schema evolution is challenging
- ii. Disadvantages:
 1. Space is wasted by repeating schema information for every record
 2. allocating variable-sized records efficiently is challenging
 3. query processing is more difficult and less efficient when the structure of the data varies
- iii. JSON & XML are best described as variable format, variable size
- iv. A VARCHAR field is best described as fixed format, variable size

d. Storing records in blocks

- i. Issues related to storing records in blocks



RAID Summary

<u>Level</u>	<u>Performance</u>	<u>Protection</u>	<u>Capacity (for N disks)</u>
0	Best (parallel read/write)	Poor (lose all on 1 failure)	N
1	Good (write slower as 2x)	Good (have drive mirror)	N / 2
5	Good (must write parity block)	Good (one drive can fail)	N - 1
6	Good (must write multiple parity blocks)	Better (can have as many drives fail as dedicated to parity)	N - X (where X is # of parity drives such as 2)

31

v. List different ways for representing strings in memory.

1. Null-terminated String: last byte value of 0 indicates end
2. Byte-length String: length of string in bytes is specified in first few bytes before string starts
3. Fixed-length String: always the same size

vi. List different ways for representing date/times in memory.

1. Date representations
 - a. Integer representation: number of days passed since a given date
 - b. String representation: show a date's components as individual characters of a string (YYYYMMDD)
2. Time representations
 - a. Integer representation: number of seconds since a given time
 - b. String representation: hours, minutes, seconds, fractions (HH:MM:SS:FF)

vii. Explain the difference between fixed and variable length records.

1. Fixed-length records: all records have the same size
2. Variable-length records: all records have the same size

viii. Compare/contrast the ways of separating fields in a record.

1. *No separator*: store length of each field, so do not need a separate separator (fixed length field). Simple but wastes space within a field.
2. *Length indicator*: store a length indicator at the start of the record (for the entire record) and a size in front of each field. Wastes space for each length field and need to know length beforehand.

12. **Utilization** =

$$(\text{records/block}) * (\text{bytes/record}) / \text{blockSize}$$
13. **Transfer time/average time to retrieve (s)** =

$$\text{numBlocks} * \text{blockSize} / \text{readBandwidth}$$

III) B-Trees

a. Objectives

i. Insert and delete from a B-tree and a B+ tree

1. Inserting

- a. find leaf node where the new key belongs (it will have 1-2 keys)
- b. if key = 1 insert new key in the node in sorted order
- c. if keys = 2 insert node in sorted order (overflow)
 - i. move middle key to parent node (split node)
 - ii. if parent keys > 3 repeat node split
- d. continue until some ancestor has only 1 node or until all ancestors are full
 - i. if all ancestors are full split root node and grow tree by 1 level

2. Splitting

- a. given a node that is overflowing, split node into 2 nodes
- b. middle value gets passed to parent node
- c. repeat until a node with room for passed value is found
- d. if root node has 2 keys, split and make a new root with the middle node

3. Deleting

- a. Locate node N containing key to delete K
 - i. if K isn't found, algorithm is complete
- b. if N = interior node, find in-order successor of K and swap with K
 - i. deletion always begins at leaf node L
- c. if L contains a value in addition to K, delete K from L
 - i. no underflow: algorithm is done
 - ii. underflow is when the # of nodes is < minimum # of keys
- d. if underflow occurs, merge node with its neighboring nodes
 - i. check L siblings
 1. if sibling has max number of keys, redistribute them

- ii. else merge L with an adjacent sibling and bring down a value from L's parent
 - iii. if parent(L) has underflow, recursively merge
 - iv. if underflow occurs at root, tree shrinks a level
 - e. Think pushing the value out of a leaf and redistribute any holes left
- ii. Calculate the maximum order of a B-tree**
 - 1. one node size = $\text{keySize} \times \text{numKeys} + \text{dataPtrSize} \times \text{numKeys} + \text{childPtrSize} \times (\text{numKeys} + 1) \leq \text{block Size}$
 - 2. max order = $\text{keySize} / 2$
 - 3.
- iii. Calculate query access times using B-tree indexes**
 - 1.
- iv. Compare/contrast B-trees and B+ trees**
 - 1. B+-trees are similar to B-trees except all key values stay in the leaves of a B+-tree
 - 2. key values removed/promoted to parent nodes from leaves are copies

IV) R-Trees

a. Objectives

- i. Explain the difference between an R-tree and a B+ tree**
 - 1. R-trees can handle multidimensional data
- ii. List some types of spatial data**
 - 1. multidimensional points
 - 2. lines
 - 3. rectangles
 - 4. geometric objects
- iii. List some types of spatial queries**
 - 1. Spatial range queries
 - a. query has associated region and asks to find matches within that region
 - b. answer may include overlapping/contained regions
 - 2. Nearest neighbor queries
 - a. find closest region to a given region
 - b. results are ordered by proximity from given region
 - 3. Spatial join queries
 - a. join two types of regions
 - b. expensive to compute
 - c. involves regions & proximity
- iv. List some applications of spatial data and queries**
 - 1. Geographic information systems (GIS)
 - a. use spatial data for modeling terrain

2. Computer-aided design (CAD)
 - a. process spatial objects when designing systems
 - b. spatial constraints
3. Multimedia databases
 - a. storing images, text, and video requires spatial data management

v. Explain the idea of insertion in an R-tree

1. start at root and go down to best fit leaf L
 - a. best fit L: child whose box needs least enlargement to cover B, resolve ties by going to smallest area child
2. if best fit L has space insert and stop
3. else split L into L1 and L2
 - a. existing entries in L + newly inserted entry must be distributed between L1 and L2

V) Hashing

a. Objectives

- i. Perform open address hashing with linear probing
 1. computes hash function($y=f(x)$) and attempts to put key in location y
 2. if y is occupied, scan array to find next open location
 - a. array is treated as circular

ii. Perform linear hashing

1. insert record with key K by computing its hash value H
2. take the last d bits of H where d is the current # of bits used
3. find the bucket m where K would belong using d
4. if $m < n$, bucket exists, go back to that bucket
 - a. insert K if bucket has space
 - b. else use an overflow block
5. if $m \geq n$, put K in bucket $m - 2^{(d-1)}$
6. after each insert, check to see if load factor $lf < \text{threshold}$
7. if $lf \geq \text{threshold}$, perform split
 - a. add new bucket n (may increase directory size d)
 - b. divide records between new bucket $n = 1bd$ and bucket $0bd$
 - c. bucket split may not be the bucket where the record was added
 - d. update n and d to reflect new bucket

iii. Define:

1. Hashing: a technique for mapping key values to locations

2. Collision: when two different keys are trying to be stored in the same location
3. Perfect hash function: a function that doesn't allow for any two keys to map to the same location
- iv. Calculate load factor of a hash table**
 1. load factor = (# of records stored)/(# of possible storage locations)
 2. # of possible storage locations (s) = numBlocks*(records/block)
- v. Compare/contrast external hashing and main memory hashing**
 1. external hashing allocates records with keys to blocks on disk rather than locations in memory
- vi. Compare/contrast B+ trees and linear hashing**
 1. hashing is better at retrieving records with a specified value for the key
 2. B+-trees are better for range queries

VI) SQL Indexing

a. Objectives

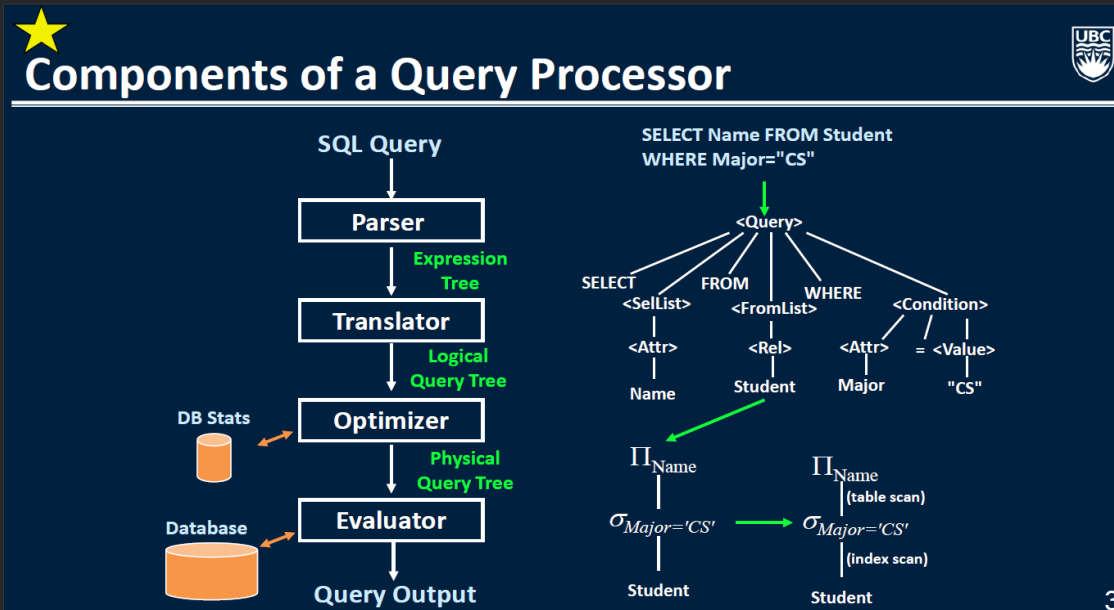
- i. Use index structures in SQL using CREATE/ALTER commands**
 1. CREATE [UNIQUE] INDEX indexName ON tableName (colName [ASC|DESC][,...])
 - a. UNIQUE: each value in index is unique
 - b. ASC/DESC: specifies sort order for index
 - c. syntax varies between systems
 2. DROP INDEX indexName
- ii. Perform insertions and searches using partitioned hashing**
 1. partitioned hashing: the overall hash location is a combination of the hash values from each key
 2. key1 ->hash function-> hash1 (h1), key2 ->hash function-> hash2 (h2)
 - a. hash location (L) = 12 bits long
 - b. 1st 6 bits = h1
 - c. 2nd 6 bits = h2
- iii. Perform searches using grid files (where $lo < x < hi$ and $lo < y < hi$)**
 1. find number of partitions for each variable
 2. number of partitions from x * number from y
- iv. Understand how bitmap indexes are used for searching and why they provide a space and speed improvement in certain cases**
 1. useful for indexing attributes that have small number values

VII) Query Processing

a. Variable definitions

- i. M-the number of buffer blocks available to the algorithm

1. always less than the size of memory
 - ii. $B(R)$ —the number of blocks on disk used to store all the tuples of R
 1. assume that R is clustered and algorithm can only read 1 block at time
 2. ignore free space in blocks
 - iii. $T(R)$ —the number of tuples in R
 - iv. $V(R,a)$ —the number of distinct values of column a in R
- b. Objectives**
- i. Diagram the query processor components and explain their function



- ii. **Calculate block access for one-pass algorithms**
 1. number of blocks accessed = $B(R) + B(S)$
 2. one pass can be completed as long as $B(S) \leq M - 1 \mid B(S) < M$
 3. where B =buffer, M =memory
- iii. **Calculate block accesses for tuple and block nested joins**
 1. R =bigger relation, S =smaller relation
 2. **Tuple Based:**
 - a. worst case = $T(R) \times T(S)$
 - b. if there's an index on the join attribute of R , the entire relation R does not have to be read
 - i. $T(S)$
 3. **Block-Based:**
 - a. worst case = $B(S) + B(R) \times \lceil \log(B(S)/(M-1)) \rceil$
 - b. if R is in outer loop = $B(S) + B(R) \times \lceil \log(B(R)/(M-1)) \rceil$
 - c. buffers smaller relation into memory
 - d.
- iv. **Perform two-pass sorting methods including all operators, sort-join and sort-merge-join and calculate performance**

- v. Perform two-pass hashing methods including all operators, hash-join and hybrid-hash-join and calculate performance
- vi. Explain the goal of query processing
 - 1. return the answer to a SQL query in the most efficient way possible given the organization of the database
- vii. List the relational and set operators
 - 1. Relational:

• selection	σ	- return subset of rows
• projection	π	- return subset of columns
• Cartesian product	\times	- all combinations of two relations
• join	\bowtie	- combines σ and \times
• duplicate elimination	δ	- eliminates duplicates

2. Set:

- Union \cup - tuple in output if in either or both
- Difference $-$ - tuple in output if in 1st but not 2nd
- Intersection \cap - tuple in output if in both
- Union compatibility means relations must have the same number of columns with compatible domains.

- viii. Explain how index and table scans work and calculate the block operations performed
 - 1. table scan: read the relation R from disk one block at a time
 - a. cost = B
 - 2. index scan: read R or just its tuples that satisfy a given condition by using an index on R
- ix. Write an iterator in java for a relational operator
 - 1. table scan:

```

init() {
    b = the first block of R;
    t = first tuple of R;
}
next() {
    if (t is past the last tuple on block b) {
        increment b to the next block;
        if (there is no next block)
            return NULL;
        else /* b is a new block */
            t = first tuple on block b;
    }
    oldt = t;
    increment t to the next tuple of b;
    return oldt;
}
close() {}

```

2. main-memory:

```

init() {
    Allocate buffer array A
    read entire relation R block-by-block into A;
    sort A using quick sort;
    tLoc = 0; // First tuple location in A
}

next() {
    if (tLoc >= T)
        return NULL;
    else
    {
        tLoc++;
        return A[tLoc-1];
    }
}

close() {}

```

How is this it
the table scan

- x. List the tuple-at-a-time relational operators
- xi. Illustrate how one-pass algorithms for selection, project, duplicate elimination, and binary operators work and calculate performance and memory requirements
- xii. Perform and calculate performance of two-pass sorting based algorithms, sort-merge algorithm, set operators, sort-merge-join/sort-join

Operators	Approximate M required	Disk I/Os
γ, δ	\sqrt{B}	$3 * B$
$\cup, -, \cap$	$\sqrt{B(R) + B(S)}$	$3 * (B(R) + B(S))$
\bowtie (sort)	$\sqrt{\max(B(R), B(S))}$	$5 * (B(R) + B(S))$
\bowtie (sort-merge)	$\sqrt{B(R) + B(S)}$	$3 * (B(R) + B(S))$

xiii. Perform and calculate performance of two-pass hashing based algorithms, hash partitioning, operation implementation and performance, hash join, hybrid-hash-join

Operators	Approximate M required	Disk I/Os
γ, δ	\sqrt{B}	$3 * B$
$\cup, -, \cap$	$\sqrt{B(S)}$	$3 * (B(R) + B(S))$
\bowtie (simple)	$\sqrt{B(S)}$	$3 * (B(R) + B(S))$
\bowtie (hybrid)	$\sqrt{B(S)}$	$(3 - \frac{2M}{B(S)})(B(R) + B(S))$

xiv. Compare/contrast sorting versus hashing methods

1. hash-based algorithms for binary operations require memory based on the size of the smaller of the two relations rather than the sum of the relation sizes
2. sort-based algorithms produce the result in sorted order which may be used for later operations
3. hash-based algorithms depend on the buckets being equal size
4. sort-based algorithms may write sorted sublists to consecutive disk blocks
5. both algorithms save disk access time by writing/reading several blocks at once if memory is available

6. hash-based joins are usually the best if neither of the input relations are sorted or there are no indexes for equi-join
7. hashing performs divide and conquer
8. sorting performs conquer and merge (sort merge)
- xv. Calculate performance of index-based algorithms**
 1. cost estimate
 2. complicated sections
 3. index joins
 - a. $\text{cost} = T(S) * (T(R) / V(R, Y))$
 - b. makes sense when $V(R, Y)$ is large and S is small
- xvi. Explain how two-pass algorithms are extended to multi-pass algorithms**
 1. two-pass algorithms based on sorting and hashing can be extended to any number of passes using recursion
 2. each pass partitions the relations into smaller pieces
 3. eventually the partitions will fit entirely into memory
 4. for k passes...
 - a. memory requirement $M = (B(R))^{1/k}$
 - b. Maximum relation size $B(R) \leq M^k$
 - c. Disk operations = $2 * k * B(R)$
 - d. w/o final pass = $2 * k * B(R) - B(R)$
- xvii. List some recent join algorithms**

VIII) Query Optimization **Not on Midterm**

a. Objectives

- i. Convert an SQL query to a parse tree using a grammar
- ii. Convert a parse tree to a logical query tree
- iii. Use heuristic optimization and relational algebra laws to optimize logical query trees
- iv. Convert a logical query tree to a physical query tree
- v. Calculate size estimates for selection, projection, joins, and set operations
- vi. Explain the difference between syntax and semantic validation and the query processor component responsible for each
- vii. Define:
 1. Valid parse tree
 2. Logical query tree
 3. Physical query tree
 4. Join-orders
 5. Left-deep
 6. Right-deep
 7. Balanced join trees

